

## Support Vector Machines

### Exercise T9.1: Structural Risk Minimization

(tutorial)

- (a) Discuss the concept of the *margin* for the linear connectionist neuron: what is the effect of a small vs. a big margin on generalization (and VC dimension)?
- (b) Write down and explain the *primal optimization problem* of model selection through structural risk minimization (SRM).
- (c) Write down the Lagrangian of the primal problem and explain the intuition behind the theorem of Kuhn and Tucker. Why can we expect sparse dual variables?
- (d) Discuss SVM classification of non-separable classes. How can this be regularized? Write down the primal problem of the C-SVM.
- (e) What is the kernel-trick and how can we exploit it?
- (f) How can multi-class problems be solved with the C-SVM method?
- (c) The Lagrangian of the SRM optimization problem for the Support Vector Machine is

$$L(\underline{\mathbf{w}}, b, \lambda_1, \dots, \lambda_p) := \frac{1}{2} \|\underline{\mathbf{w}}\|^2 - \sum_{\alpha=1}^p \lambda_{\alpha} \left\{ y_T^{(\alpha)} (\underline{\mathbf{w}}^T \underline{\mathbf{x}}^{(\alpha)} + b) - 1 \right\}.$$

The derivative w.r.t.  $\underline{\mathbf{w}}$  and  $b$  are:

$$\frac{\partial L}{\partial \underline{\mathbf{w}}} = \underline{\mathbf{w}} - \sum_{\alpha=1}^p \lambda_{\alpha} y_T^{(\alpha)} \underline{\mathbf{x}}^{(\alpha)} \stackrel{!}{=} 0 \quad \text{and} \quad \frac{\partial L}{\partial b} = - \sum_{\alpha=1}^p \lambda_{\alpha} y_T^{(\alpha)} \stackrel{!}{=} 0.$$

Substituting  $\underline{\mathbf{w}} = \sum_{\alpha=1}^p \lambda_{\alpha} y_T^{(\alpha)} \underline{\mathbf{x}}^{(\alpha)}$  into the original Lagrangian yields the dual:

$$\max_{\lambda} L(\lambda_1, \dots, \lambda_p) := -\frac{1}{2} \sum_{\alpha, \beta=1}^p \lambda_{\alpha} \lambda_{\beta} y_T^{(\alpha)} y_T^{(\beta)} \underline{\mathbf{x}}^{(\alpha)\top} \underline{\mathbf{x}}^{(\beta)} + \sum_{\alpha=1}^p \lambda_{\alpha}$$

$$\text{s.t.} \quad \lambda_{\alpha} \geq 0, \quad \forall \alpha \quad \text{and} \quad \sum_{\alpha=1}^p \lambda_{\alpha} y_T^{(\alpha)} = 0.$$

- (d) C-SVM defines slack-variables  $\varphi_{\alpha} \geq 0$  for all samples, and the primal problem is

$$\min_{\underline{\mathbf{w}}, b} \frac{1}{2} \|\underline{\mathbf{w}}\|^2 + \frac{C}{p} \sum_{\alpha=1}^p \varphi_{\alpha} \quad \text{s.t.} \quad y_T^{(\alpha)} (\underline{\mathbf{w}}^T \underline{\mathbf{x}}^{(\alpha)} + b) \geq 1 - \varphi_{\alpha}, \quad \text{and} \quad \varphi_{\alpha} \geq 0, \forall \alpha.$$

### Exercise H9.1: Deriving the C-SVM optimization problem (homework, 3 points)

- (a) (1 point) Linear connectionist neurons have a degree of freedom that is not used in classification. By setting the constraint

$$\min_{\alpha=1,\dots,p} \left| \underline{\mathbf{w}}^\top \underline{\mathbf{x}}^{(\alpha)} + b \right| \stackrel{!}{=} 1$$

this degree is eliminated. Show that under this constraint the Euclidean distance  $d(\underline{\mathbf{x}}^{(\alpha)}, \underline{\mathbf{w}}, b)$  of sample  $\underline{\mathbf{x}}^{(\alpha)}$  to the closest point of the decision boundary  $\{x|y(x) = 0\}$  is bounded by

$$d(\underline{\mathbf{x}}^{(\alpha)}, \underline{\mathbf{w}}, b) \geq \frac{1}{\|\underline{\mathbf{w}}\|}, \quad \forall \alpha \in \{1, \dots, p\}.$$

- (b) (2 points) Write down the Lagrangian of the primal optimization problem of the C-SVM and derive the dual optimization problem of the C-SVM:

$$\max_{\lambda} \left\{ -\frac{1}{2} \sum_{\alpha=1}^p \sum_{\beta=1}^p \lambda_{\alpha} \lambda_{\beta} y_T^{(\alpha)} y_T^{(\beta)} \left( \mathbf{x}^{(\alpha)} \right)^\top \mathbf{x}^{(\beta)} + \sum_{\alpha=1}^p \lambda_{\alpha} \right\}$$

with  $0 \leq \lambda_{\alpha} \leq \frac{C}{p}, \forall \alpha, \quad \text{and} \quad \sum_{\alpha=1}^p \lambda_{\alpha} y_T^{(\alpha)} = 0.$

### Exercise H9.2: C-SVM with standard parameters (homework, 3 points)

In this exercise, we use C-SVMs to solve the “XOR”-classification problem from exercise sheet 7. To this end (1) first create a *training set* of 80 data as described in exercise H7.1 and (2) create a *test set* of 80 data from the same distribution.<sup>1</sup>

You can use existing software: `libsvm`<sup>2</sup> implements optimization routines (Matlab & Python) for SVMs. Alternatively, you can use the corresponding `scikit.learn` class<sup>3</sup>. For R, the package `e1071` implements SVM-optimization.

- Download, install, and familiarize yourself with `LIBSVM` or one of the other packages.
- Read the *Practical Guide to Support Vector Classification*<sup>4</sup> especially section 3.2 on *Cross-Validation*.

Next, use your chosen SVM implementation to train a C-SVM with RBF kernel and the software’s standard parameters. Classify the test data and report the classification error quantified by the 0/1 loss function (percentage of wrong predictions). Visualize the results as in exercise H7.2: plot the training patterns and the decision boundary (e.g. with a contour plot) in input space. Additionally, highlight the support vectors.

### Exercise H9.3: C-SVM parameter optimization (homework, 4 points)

<sup>1</sup>Since for SVMs typically labels  $y_T \in \{-1, +1\}$  are used instead of  $y_T \in \{0, 1\}$  it might be beneficial to use the former and not the latter.

<sup>2</sup><http://www.csie.ntu.edu.tw/~cjlin/libsvm/>

<sup>3</sup><http://tinyurl.com/lrpxw9k>

<sup>4</sup><http://www.csie.ntu.edu.tw/~cjlin/papers/guide/guide.pdf>

- (a) (2 points) Use cross-validation and grid-search to determine good values<sup>5</sup> for the  $C$  and the kernel parameter  $\gamma$ . Follow the procedure described in the *guide*: Define the grid using exponentially growing sequences of  $C$  and  $\gamma$ , e.g.  $C \in \{2^{-6}, 2^{-4}, \dots, 2^{10}\}$ ,  $\gamma \in \{2^{-5}, 2^{-3}, \dots, 2^9\}$ . Make sure you only use the training data in this step. Plot the mean training-set classification rate and cross-validation performance as a function of  $C$  and  $\gamma$  (e.g. using contour plots as in figure 2 of the *guide*).
- (b) (1 point) Find the best combination of  $C$  and  $\gamma$  and train the RBF C-SVM on the *entire* training data, this time using these “optimal” parameters. Plot the results in the same way as in exercise H9.2.
- (c) (1 point) Compare the results with those obtained in H9.2, both in terms of statistics (e.g. classification performance, number of support vectors) and visually (e.g. signs of over- and under-fitting). What happens when you divide  $C$  or  $\gamma$  by 4?

**Total 10 points.**

---

<sup>5</sup>Note that  $\gamma = 1/(2\sigma^2)$  in the lecture notation and that the  $C$  from the guide corresponds to  $C/p$  from the lecture; you do not have to use the lecture’s notation but can stick to the (more standard) variant from the guide.